

CSDS 313/413: Introduction to Data Analysis

Homework 3: Pairwise Association

Solutions

Wiam Skakri and Elizabeth Kitakule

11/10/2025

1 Task 1: Association Between Genomic Variants

1.1 Part (a): Two Variables from p1a.csv

1.1.1 Test Statistics and P-values

Mutual Information (MI)

- Observed MI value: 0.046993
- Number of permutations (N): 10,000
- P-value: 0.0027
- Selected significance level (α): 0.05
- Conclusion: Statistically significant association (p-value < α)

Jaccard Index (JI)

- Observed JI value: 0.000000
- Number of permutations (N): 10,000
- P-value: 1.0000
- Selected significance level (α): 0.05
- Conclusion: No statistically significant association (p-value $\geq \alpha$)

Pearson's Chi-squared (χ^2)

- Observed χ^2 value: 6.455974
- Degrees of freedom: 1
- P-value: 0.01106
- Selected significance level (α): 0.05
- Conclusion: Statistically significant association (p-value < α)

1.1.2 Analysis and Interpretation

Statistical Significance: The three statistics show conflicting results regarding statistical significance at the $\alpha = 0.05$ level. Two out of three statistics (Mutual Information with p-value = 0.0027 and Pearson's χ^2 with p-value = 0.01106) indicate a statistically significant association between the two genomic variants. However, the Jaccard Index shows no significant association (p-value = 1.0000). This discrepancy suggests that while there is some statistical dependence between the variables, it may not manifest as co-occurrence of positive cases (both variants present simultaneously).

Strength of Association: The strength of association appears to be weak based on the observed test statistics:

- **Mutual Information (0.046993):** MI ranges from 0 (complete independence) to 1 for binary variables. The observed value of 0.047 is very close to 0, indicating a very weak association. This suggests that knowing the state of one genomic variant provides minimal information about the other.
- **Jaccard Index (0.000000):** JI ranges from 0 (no overlap) to 1 (complete overlap). A value of exactly 0 indicates there is no co-occurrence of both variants being present (both equal to 1) simultaneously in any individual. Looking at the contingency table, we see that when variable X = 1, variable Y is always 0 (50 cases), confirming zero overlap.
- **Pearson's χ^2 (6.456):** The magnitude itself is difficult to interpret without context, but the relatively low p-value (0.01106) indicates the departure from independence is unlikely to be due to chance. The effect size, however, appears modest given the sample size of 199 individuals.

Agreement Between Statistics: The statistics show partial agreement: 2 out of 3 tests (MI and χ^2) detect a statistically significant association at $\alpha = 0.05$, while the Jaccard Index does not. This partial disagreement is informative about the nature of the association present in the data.

Explanation of Any Discrepancies: The discrepancy between statistics can be explained by their different mathematical properties and what they measure:

- **Mutual Information** captures any form of statistical dependence between variables using an information-theoretic approach. It detected significance because there is a deviation from independence in the joint distribution—specifically, when X = 1, Y is never 1 (Y is always 0). This negative association (mutual exclusivity pattern) is captured by MI.
- **Jaccard Index** specifically measures co-occurrence of positive cases (overlap when both variables equal 1). Since the contingency table shows that there are zero cases where both X = 1 and Y = 1 simultaneously, JI = 0 and shows no significance. JI is not sensitive to the negative association pattern present in this data.
- **Pearson's χ^2 test** evaluates whether the observed frequencies in the contingency table differ significantly from what would be expected under independence. The test detected significance

because the observed cell count of 0 (when $X = 1$ and $Y = 1$) differs notably from the expected count of approximately 5.28 under independence.

In summary, the data exhibits a pattern of negative association or mutual exclusivity—when one variant is present, the other tends to be absent. This pattern is captured by MI and χ^2 but not by JI, which only measures positive co-occurrence. The disagreement among statistics is not a contradiction but rather reveals different aspects of the association structure in the data.

1.1.3 Visualization




Figure 1: Permutation test results and chi-squared distribution for Task 1(a). Left: Mutual Information permutation test showing the null distribution (blue) and observed value (red dashed line). Center: Jaccard Index permutation test. Right: Pearson's chi-squared test showing the theoretical χ^2 distribution and the observed test statistic. The observed MI and χ^2 values fall in the tail of their respective distributions, indicating statistical significance.

1.2 Part (b): 105 Variable Pairs from p1b.csv

1.2.1 Test Statistics and P-values

Number of Permutations (N): 10,000

Selected Significance Level (α): 0.05

Multiple Hypothesis Correction Method: Benjamini-Hochberg (BH) procedure

1.2.2 Results Summary

Mutual Information (MI)

- Number of significantly associated pairs: 91 out of 105 (86.7%)

Jaccard Index (JI)

- Number of significantly associated pairs: 56 out of 105 (53.3%)

Pearson's Chi-squared (χ^2)

- Number of significantly associated pairs: 90 out of 105 (85.7%)

1.2.3 Comparison of Statistics

Overlap Between Statistics:

- MI and JI overlap: 55 pairs (60.4% of MI significant, 98.2% of JI significant)
- MI and χ^2 overlap: 89 pairs (97.8% of MI significant, 98.9% of χ^2 significant)
- JI and χ^2 overlap: 56 pairs (100% of JI significant, 62.2% of χ^2 significant)
- All three statistics overlap: 55 pairs (52.4% of all 105 pairs)

The overlap analysis reveals that nearly all pairs identified as significant by the Jaccard Index are also identified by the other two statistics. However, MI and χ^2 identify additional significant pairs that JI does not detect.

Which Two Statistics Are Most Similar: Based on Spearman correlation analysis of the test statistic values across all 105 pairs:

- MI vs JI correlation: $\rho = 0.683$
- MI vs χ^2 correlation: $\rho = 0.989$
- JI vs χ^2 correlation: $\rho = 0.737$

Mutual Information and χ^2 are the most similar pair ($\rho = 0.989$), showing nearly perfect rank correlation. This makes sense because both statistics measure general statistical dependence between variables, whereas JI specifically measures co-occurrence of positive cases.

Preferred Test Statistic and Justification: For genomic variant association studies, **I would prefer Mutual Information (MI)** for the following reasons:

1. **Captures all types of dependence:** MI detects both positive associations (co-occurrence) and negative associations (mutual exclusivity), as demonstrated in Part (a). This is crucial in genomics where variants may be mutually exclusive due to biological constraints.
2. **Information-theoretic interpretation:** MI directly quantifies how much information one variant provides about another, which has clear biological meaning.
3. **Non-parametric:** Unlike χ^2 , MI makes no distributional assumptions and relies on permutation testing, making it robust for sparse data common in genomics.
4. **Comprehensive detection:** MI identified 91 significant pairs, capturing nearly all pairs found by χ^2 (89 overlap) and JI (55 overlap), plus additional associations.

Impact of Using Only Preferred Statistic: If we used only Mutual Information among the three statistics tested:

- **Minimal loss of information:** We would identify 91 significant pairs, missing only 1 pair that χ^2 uniquely identified (V1-V9) and 1 pair that JI uniquely identified (none that weren't also in MI).
- **More comprehensive coverage:** We would capture 2 additional pairs that χ^2 missed (V4-V9) and 36 additional pairs that JI missed, suggesting MI has higher sensitivity while maintaining good specificity after BH correction.
- **Biological interpretation:** We would not lose significant biological insights. The high correlation between MI and χ^2 ($\rho = 0.989$) suggests they provide largely redundant information, while MI's ability to detect associations missed by JI means we would gain rather than lose information by using MI alone.

1.2.4 Visualizations




Figure 2: Comprehensive comparison of MI, JI, and χ^2 for 105 variable pairs. **Top row:** Scatter plots showing the relationship between test statistics, with significant pairs marked. The near-linear relationship between MI and χ^2 (center) confirms their high correlation ($\rho = 0.989$). **Middle row:** P-value distributions for each statistic, showing enrichment of small p-values indicating true associations. **Bottom:** Overlap analysis showing the number of pairs uniquely or jointly identified as significant by each statistic combination. The visualization demonstrates that MI and χ^2 show strong agreement, while JI identifies a subset of associations focused on co-occurrence patterns.

2 Task 2: Association Between Continuous Variables

2.1 Part (a): Variable Pair from p2a.csv

2.1.1 Test Statistics and P-value

Pearson Correlation (r_a):

- Correlation coefficient: 0.38087503578373005
- P-value (p_a): 1.0409455130062156e-83
- Selected significance level (α): 0.05
- Sample size: 2400 samples

2.1.2 Analysis and Interpretation

Statistical Significance: The p-value is an extremely small value, 1.04 e-83. This is far below the chosen significance level, 0.05. Therefore, we reject the null hypothesis of no linear relationship — there is a statistically significant positive linear relationship between the two variables in p2a. The correlation likely did not happen by random chance.

Magnitude of Association: The correlation coefficient, 0.38, indicates a moderate association between the two variables. There is a relationship present but it is not a strong one. The variation in variable 1 is not wholly explained by the other variable in the dataset.

Direction of Association: The correlation is positive. Given that the correlation coefficient has a positive sign, it indicates that as one variable increases, the other increases. The association indicates an upward trend in the sample size.

2.1.3 Visualization

Include scatter plot for p2a.csv




Figure 3: Scatter plot for p2a.csv

2.2 Part (b): Comparison of p2a.csv and p2b.csv

2.2.1 Test Statistics and P-value for p2b.csv

Pearson Correlation (r_b):

- Correlation coefficient: 0.9312196333264214
- P-value (p_b): 3.73732100843861e-49
- Selected significance level (α): 0.05
- Sample size: 110 samples

2.2.2 Comparison Analysis

Comparison of Correlations (r_a vs r_b): r_b indicates a stronger relation than r_a . r_b is 0.931, which indicates a strong association between the two variables in p2b.csv. r_a showed us a moderate association that displayed an association but not at the same magnitude of r_b .

Stronger Association Based on Correlations: r_b

Comparison of P-values (p_a vs p_b): p_b is extremely small and far below the significance level, 0.05. The p-value, p_a , is also extremely small and both indicate the statistical significance of the association between the two variables in each dataset. We reject the null hypothesis that there is no association between the variables. Including the strong r_b , the values show a clear and statistically significant linear relationship between the variables.

Stronger Association Based on P-values: p_a has a stronger association based on p-value even though it has a lower correlation coefficient.

Agreement Between Correlations and P-values: No. Judgement based on solely correlation coefficients indicate that the data set p2b has a stronger association. However, judgement based solely on p-values, p2a has a stronger association based off the smaller p-value.

Explanation of Any Discrepancy: The difference in p-value could be due to difference in sample sizes. A large dataset like that of p2a will result in smaller p-values regardless of the fact that there is a weaker association. The correlation coefficient directly reports the strength of the association, regardless of the dataset size.

Visual Assessment from Scatter Plots: The scatterplots display the difference in sample size between p2a and p2b. Therefore, adding to the p-value discrepancy and its justification. Both the scatterplots show an upward or positive linear association.

2.2.3 Visualizations




Figure 4: Scatter plots for p2a.csv (left) and p2b.csv (right)

2.3 Part (c): Comparison of p2a.csv and p2c.csv

2.3.1 Test Statistics and P-value for p2c.csv

Pearson Correlation (r_c):

- Correlation coefficient: 0.0411789977768318
- P-value (p_c): 0.059195916605514404
- Selected significance level (α): 0.05
- Sample size: 2100 samples

2.3.2 Comparison Analysis

Comparison of Correlations (r_a vs r_c): r_c indicates a weaker relation than r_a . r_c is 0.0412, which indicates a weak association between the two variables in p2c. r_a showed us a moderate association that displayed an association at a higher magnitude than r_c . r_c shows little association between the two variables in the dataset.

Stronger Association Based on Correlations: r_a has a stronger association with 0.381.

Comparison of P-values (p_a vs p_c): p_c is greater than the chosen significance level, 0.05. p_c is 0.0591. In comparison, p_a is significantly smaller than the significance level. therefore, we fail to reject the null hypothesis that there is no association between the variables in the dataset, p2c. p_c indicates that the relationship is not statistically significant. On the other hand, The p-value, p_a , is extremely small and indicates the statistical significance of the association between the two variables in each dataset. We reject the null hypothesis that there is no association between the variables

Stronger Association Based on P-values: p_a indicates a stronger association in the dataset, p2a.

Agreement Between Correlations and P-values: Yes. The correlation coefficients and p-values allow for the same conclusion when comparing the linear relationship between the two datasets. Both the p-value and the correlation indicate a linear relationship, however moderate , that exists in dataset, p2a. The same values for dataset, p2c, indicate a lack of a linear relationship and display a weak or no association between the variables in the dataset.

Visual Assessment from Scatter Plots: The scatterplot of the variables in p2c display the lack of linear association between the two variables. Unlike the scatter plots for p2a and p2c where there was a clear upward trend with an increase in each variable, here, the diagram shows that there is a decrease and increase of the variables. The change in variable 1 has no predictable linear effect on variable 2.

Agreement with Statistical Measures: The visual assessment agrees with the statistical measures as there is no visible linear association.

2.3.3 Visualizations




Figure 5: Scatter plots for p2a.csv (left) and p2c.csv (right)

3 Code Appendix